



RUTGERS

# The Next Hit

Predicting Movie Success and Building A Recommendation Engine

Capstone Project for Massive Data Mining - Spring '18

## Team 14

Tahiya Chowdhury, George Chantzialexiou, Christos Mitropoulos

# Motivation of the project

- Modern video-on-demand companies like Netflix, Amazon, Youtube RED have recently started **producing original content**.
- These companies need to **keep users satisfied** by making good movie recommendations and **maximize their profit** by producing original content.
- There is a need for **recommendation systems, movie revenue predictions** and **prediction of a movie's success/failure** to provide streamlined experience for both users and the companies

# Contribution

1. **Movie-revenue prediction** system to help companies make better investments.
2. **Binary classifier** that predicts if a movie will make a profit or a loss
3. Two types of movie-recommendation systems:
  - a. Traditional **recommendation system**
  - b. **Online - streaming** recommendation system

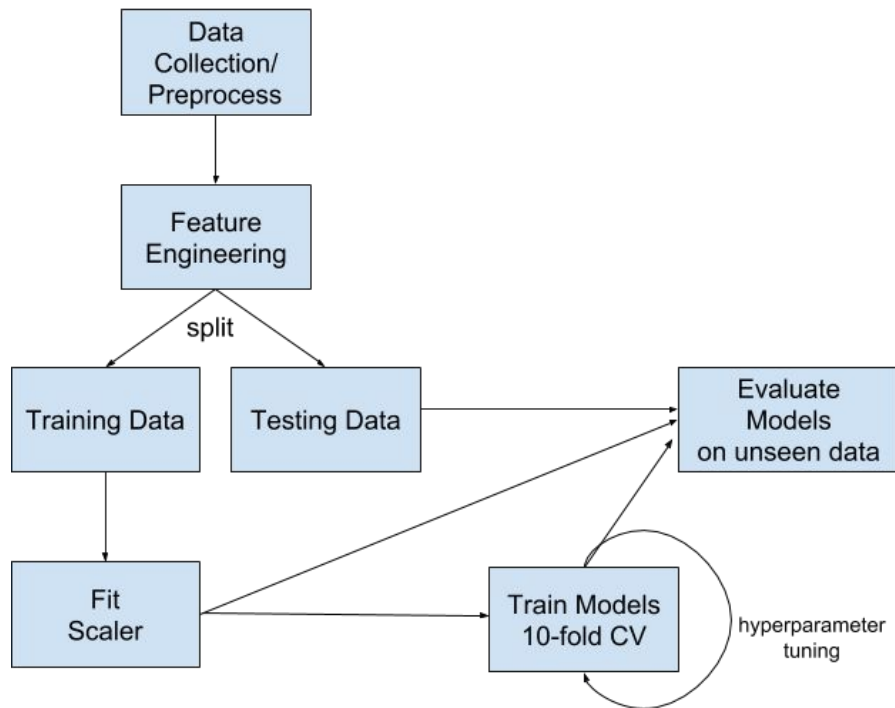
# Dataset Description

- An ensemble of data collected from **TMDB** and **GroupLens**. The Movie Details, Credits and Keywords have been collected from the **TMDB Open API** (on or before July 2017.)
- These files contain:
  - Metadata for all 45,000 movies listed in the Full MovieLens Dataset.  
Data points include : **cast, crew, plot keywords, budget, revenue, posters, release dates, languages, production companies, countries, TMDB vote counts and vote averages.**
  - Files containing 26 million ratings from 270,000 users for all 45,000 movies.  
Ratings are on a scale of 1-5 and have been obtained from the official GroupLens website.
- There are 5000 movies for which we have data on revenue and budget ratio.

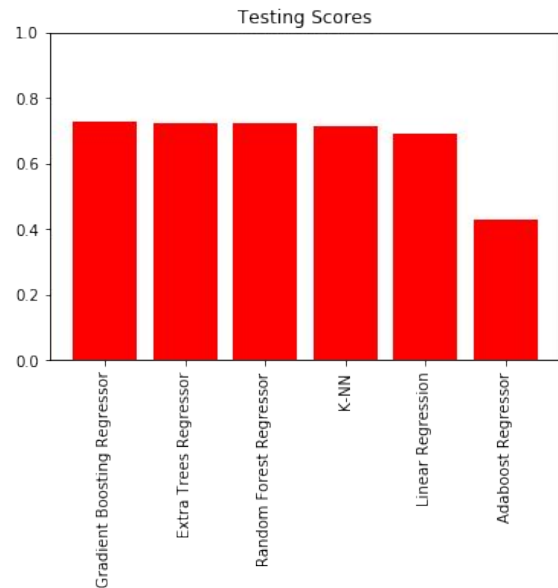
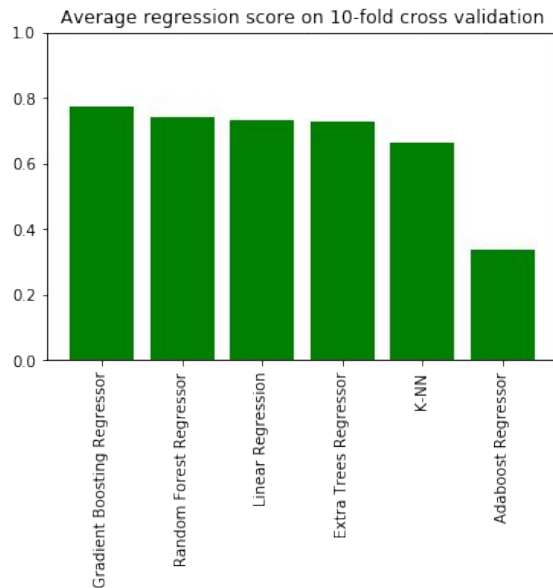
# Movie revenue prediction - Binary Classification

Goal: Facilitate production companies in decision making by building

1. Revenue prediction tool.
2. Classification tool to predict if a movie will make profit or loss.







# Results of Movie Revenue Prediction



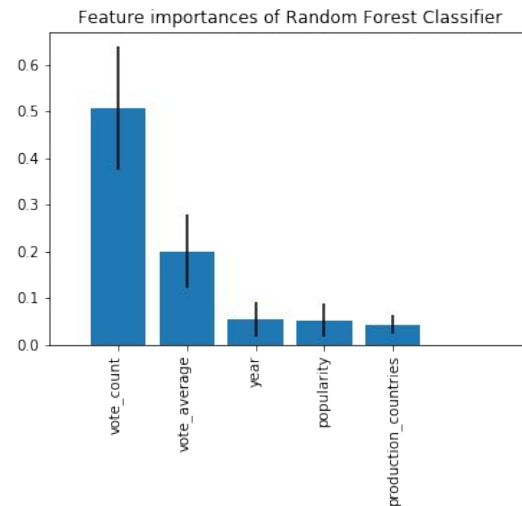
Best performing model is the Gradient Boosting Regressor (CV-score: 0.77, Test Score: 0.72 ).  
Most models, except AdaBoost Regressor, have similar performance.

Note: Score reported is *the coefficient of determination  $R^2$  of the prediction*.

# Results of Binary Classification into profit or loss



- **Decision Trees** have higher **accuracy** in the Test set, however the **Random Forest Classifier** performs better with regards to **recall**.
- A balanced dataset could provide better results.
- Feature importance can provide interesting conclusions



# Recommender System

## Content-based Recommender

- Useful when a new user enters the system and has no ratings (**cold start problem**)
- Recommends movies based on particular movie or genre

## Our approach

- **Weighted vote, vote count, vote average, genres** from the movie metadata
- **Top 3 cast members** and **director** from cast data
- Frequently appeared **keywords** from keywords data
- Computed Cosine similarity score of the resultant matrix (feature extraction)
- Recommended movies based on user suggested item



# Recommendation result for a particular movie

```
recommendations('The Godfather')
```

	title	year	vote_count	vote_average	genres	wr
1199	The Godfather: Part II	1974	3418	8	[Drama, Crime]	7.689586
1186	Apocalypse Now	1979	2112	8	[Drama, War]	7.530356
1934	The Godfather: Part III	1990	1589	7	[Crime, Drama, Thriller]	6.623473
1312	Dracula	1992	1087	7	[Romance, Horror]	6.499201
3635	The Conversation	1974	377	7	[Crime, Drama, Mystery]	6.060771
24284	The Drop	2014	859	6	[Drama, Crime]	5.746547
2025	The Outsiders	1983	293	6	[Crime, Drama]	5.549223
1614	The Rainmaker	1997	239	6	[Drama, Crime, Thriller]	5.513054
8911	Rumble Fish	1983	141	6	[Action, Adventure, Crime, Drama, Romance]	5.430061
754	Jack	1996	340	5	[Comedy, Drama, Science Fiction]	5.137319

# Collaborative Filtering based recommender

Useful for recommending movies by predicting ratings on unseen movies by a user

## Our Approach

- **Used Surprise library for Python**

	<b>SVD</b>	<b>SlopeOne</b>	<b>KNNBaseline</b>
<b>Mean RMSE</b>	0.8968	0.9286	0.8972
<b>Mean MAE</b>	0.6906	0.7109	0.8972

- Built a test set of movies that are unseen by the users
- Created Recommendation list for each user's unseen movies based on the predicted ratings

# Recommendation result for a particular user

	uid	iid	r_ui	est
<b>2034527</b>	4	73290	3.543608	5.000000
<b>2037288</b>	4	55732	3.543608	5.000000
<b>2034088</b>	4	44197	3.543608	5.000000
<b>2029805</b>	4	1252	3.543608	5.000000
<b>2030579</b>	4	1228	3.543608	5.000000
<b>2030055</b>	4	1221	3.543608	5.000000
<b>2030066</b>	4	912	3.543608	5.000000
<b>2029786</b>	4	608	3.543608	5.000000
<b>2030285</b>	4	318	3.543608	5.000000
<b>2029877</b>	4	50	3.543608	5.000000

## Problem

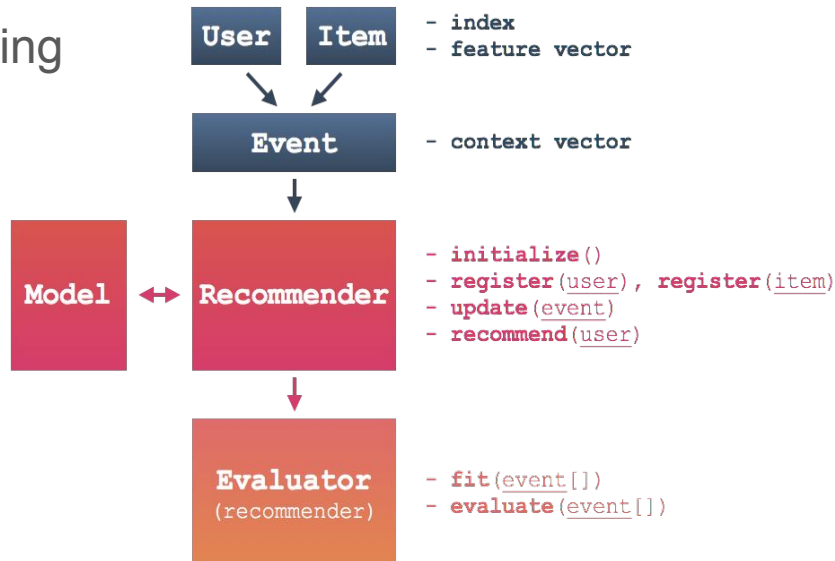
**New ratings are coming all the time!**

# Streaming Recommendation - Goal

- Traditional training method is not scalable
- Companies (Netflix) keep receiving new ratings from Users
  - New ratings should be accounted to the recommender system.
- Users tastes are changing overtime
  - System should be able to evaluate based on the latest choices of the user

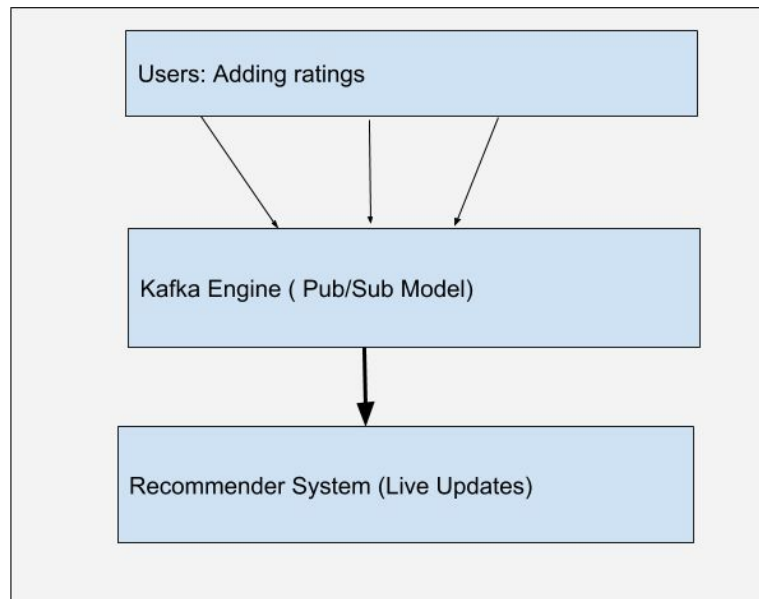
# Streaming Recommendation - How

- Online Updating via Structured streaming
- Model:
  - Collaborative: K-Nearest Neighbors





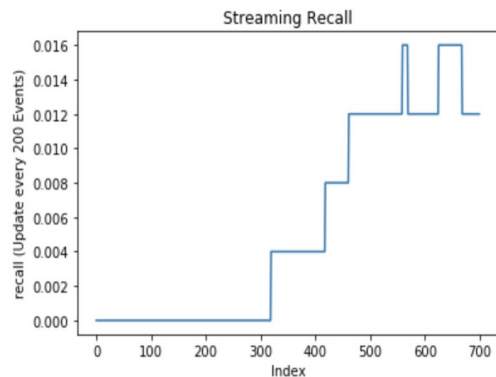
# Streaming Recommendation - How

- Streaming Architecture
  - Apache Kafka : Distributed Messaging System
  - Structure Recommender System
- Apache Kafka
  - Push ( Input from user)
  - Pull ( Recommender pulls new data and update Recommender Engine



# Streaming Recommendation - Results

- Evaluation Method: Test-and-Learn
  - Initial batch training
  - Later: Recommender sequentially launch top-10 recommendation and check if observed item is correctly included in the recommendation list
  - Continuous Monitoring of live updates



Questions?



Thank you!

# References

1. Incremental Factorization Machines for Persistently Cold-starting Online Item Recommendation ( <https://arxiv.org/pdf/1607.02858.pdf> )
2. Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011. (<http://jmlr.csail.mit.edu/papers/v12/pedregosa11a.html>)
3. Surprise: A python scikit for recommender system (<http://surpriselib.com>)